

Computer Vision

Image segmentation

November 7, 2024

1 Environment setup

In this assignment you will 1) implement the mean-shift algorithm using PyTorch and validate it on an image 2) implement a simplified version of SegNet (<https://arxiv.org/abs/1511.00561>) and train/validate your model on the provided multi-digit MNIST dataset.

Running PyTorch on CPU would suffice for this assignment and it should take around 10 minutes to train the simplified SegNet on CPUs. However, you are free to use GPU if you have access to one.

2 Mean-Shift Algorithm (Total: 40 pts)

For this task, you can obtain the skeleton code in the provided `mean-shift.zip` file. You will need to implement three placeholder functions to make the skeleton code runnable (35 pts total). Then, you will need to accelerate the naive implementation (5 pts).

2.1 Implement the distance Function (12 pts)

The `distance` function should compute the distance between a given point and all other points that are within a specific radius. For simplicity, here we consider this radius to be $+\infty$ such that you don't have to explicitly handle it.

2.2 Implement the gaussian Function (12 pts)

The `gaussian` function should compute the weights of points according to the distance computed by the `distance` function.

2.3 Implement the update_point Function (11 pts)

The `update_point` function should update point position according to weights computed from the `gaussian` function.

2.4 Accelerating the Naive Implementation (5 pts)

How is the running time of your implemented mean-shift algorithm? One way to accelerate it would be to vectorize the inputs to avoid looping over every point. Implement such vectorized version and report the speed-up you gain from such improvement. For timing, please make sure you run your experiments consistently on the same device (CPU or GPU).

Note: Usually, vectorized inputs would result in faster inference, but it also depends on your hardware (CPU/GPU speed, RAM speed and capacity, and background programs). As long as your implementation is correct and you have reasonable explanation in your report, you will get full points on this subtask.

3 Implement and Train a Simplified Version of SegNet (Total: 60 pts)

You can obtain the skeleton code in the provided `seg-net.zip` file. Please follow the instructions in `README.md` to install the package and implement a simplified version of SegNet (40 pts). You also need to train the model, and validate its accuracy on the validation set (20 pts).

3.1 Implement a Simplified Version of SegNet (40 pts)

You will need to implement simplified SegNet model in `lib/models/seg_net_lite.py`. Other modules, including dataloader, pre-/post-processing and training/evaluation scripts are provided with proper comments. Although you will not need to modify other parts of the codebase if you implement SegNet properly, it is still highly recommended to read through the codebase to get a good understanding of how the full pipeline works.

3.2 Train and Validate Your Model (20 pts)

Train and validate your model using the provided training and validation scripts. Details can be found in `README.md` of the provided codebase. For this task, you get 20 points if your final validation accuracy (mean IoU) surpasses 0.8. The points will be scaled linearly between 0-20 if your final validation accuracy is less than 0.8. Please hand-in your trained model (a *single* file named `model_best.pth.tar`) along with your modified `seg_net_lite.py` file. Note that if your `seg_net_lite.py` file cannot load your model, then you will get 0 (out of 20) point for this subtask. It is recommended to use a version control tool (e.g. Git) to keep track of your modifications, so that you can make sure your final version before hand-in does not modify anything other than `lib/models/seg_net_lite.py`.

4 Hand-in

Please hand-in 1) `mean-shift.py` 2) `seg_net_lite.py` 3) `model_best.pth.tar` 4) A short report (1 page at max.) describing your implementation of mean-shift, including the timings of for-loop-based vs. vectorized mean-shift, all in a *single* zip file. You do not need to report anything on the SegNet task, we will just check your implementation and model accuracy.